

Introduction to Random Numbers, Sampling, and MCMC Methods

Bill Watson

S&P Global

August 7, 2019

What is Sampling and why is it useful?

- ▶ Sampling is the practice of generating observations from a population
- ▶ Monte Carlo methods are algorithms that rely on repeated random sampling to obtain approximations where it is difficult or impossible to use deterministic approaches
 - ▶ Optimization
 - ▶ Numerical Integration
 - ▶ Sampling distributions

Application: Approximating π

Algorithm 1 Approximating π

Input: Batch Size N


1: Sample $u_1 \sim \text{Uniform}(0, 1)$ N times

2: Sample $u_2 \sim \text{Uniform}(0, 1)$ N times

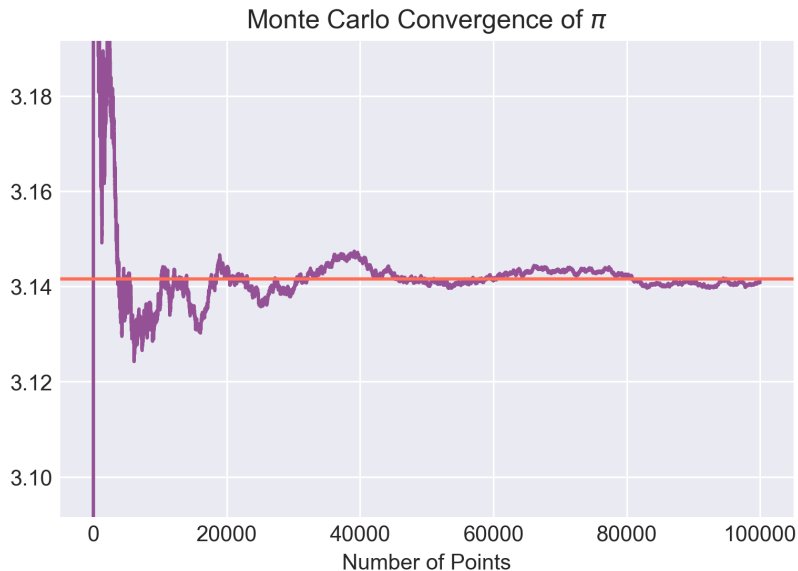
3: $\tilde{\pi} = \frac{4}{N} \cdot \left| \left\{ (u_1, u_2) \mid \sqrt{u_1^2 + u_2^2} < 1 \right\} \right|$

Output: $\tilde{\pi}$

Application: Approximating π



Application: Approximating π



What can we do with our samples?

- ▶ Integration: $\int p(x) f(x) dx = \frac{1}{n} \sum_{x_i \sim P} f(x_i)$
- ▶ Expectation: $\mu = \frac{1}{n} \sum_{x_i} x_i$
- ▶ Variance: $\sigma^2 = \frac{1}{n} \sum_{x_i} (x_i - \mu)^2$
- ▶ Median: $\text{median} = \text{median}(x_1, x_2, \dots, x_n)$
- ▶ Entropy: $\mathbb{H}(P) = -\frac{1}{n} \sum_{x_i \sim P} \log p(x_i)$
- ▶ CDF: $p(c) = \frac{1}{n} |\{x_i \mid x_i \leq c\}|$

Pseudo-Random Number Generators

Pseudo-Random Number Generators: LCG

Algorithm 2 Linear Congruential Generator

Input: Modulus m , Multiplier a , Increment c , Seed X_0

1: $X_{i+1} = (a \cdot X_i + c) \bmod m$

Output: X_{i+1}

Pseudo-Random Number Generators: Mersenne Twister

Generating a Normal from the CDF

- ▶ We can use the cumulative distribution function to sample any distribution
- ▶ For instance, a normal's CDF is:

$$F(x) = \frac{1}{2} \left[1 + \operatorname{erf} \left(\frac{x - \mu}{\sigma\sqrt{2}} \right) \right]$$

- ▶ With an Inverse CDF as:

$$F^{-1}(p) = \mu + \sigma\sqrt{2} \cdot \operatorname{erf}^{-1}(2p - 1)$$

- ▶ $\operatorname{erf}(x)$ is the error function, defined as:

$$\operatorname{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$$

Inverse Transform Sampling

- ▶ It's easy to generalize this method to any distribution with a closed-form inverse CDF

Algorithm 3 Inverse Transform Sampling

Input: Inverse CDF F^{-1}

- 1: Sample $u \sim \text{Uniform}(0, 1)$
- 2: $X = F^{-1}(u)$

Output: X

Inverse Transform Sampling: Example

Table of Inverse CDFs for Common Distributions

Distribution	$F^{-1}(p)$
$\mathcal{N}(\mu, \sigma)$	$\mu + \sigma\sqrt{2} \cdot \operatorname{erf}^{-1}(2p - 1)$
Uniform(a, b)	$a + p \cdot (b - a)$
Exp(λ)	$\frac{-\ln(1-p)}{\lambda}$

Rejection Sampling

Rejection Sampling: Algorithm


Algorithm 4 Rejection Sampling

Input: Model P , Proposal Q , $M > 1$

- 1: Sample $x \sim P$
- 2: Sample $u \sim \text{Uniform}(0, 1)$
- 3: **if** $u < \frac{p(x)}{M \cdot q(x)}$ **then**
- 4: Accept x
- 5: **else**
- 6: Reject x
- 7: **end if**

Output: Accepted Samples

Rejection Sampling: Intuition



Rejection Sampling: Example




Figure: Rejection Sampling with $M = 1.3$

Rejection Sampling: Example




Figure: $M = 1.3$, 6,688 Accepted Samples

Rejection Sampling: Example




Figure: Rejection Sampling with $M = 2.5$

Rejection Sampling: Example




Figure: $M = 2.5$, 4,085 Accepted Samples

Rejection Sampling: Example




Figure: Rejection Sampling with $M = 100$

Rejection Sampling: Example




Figure: $M = 100$, 92 Accepted Samples

Rejection Sampling: Pros & Cons

What is MCMC?

Metropolis-Hastings

Metropolis-Hastings: Algorithm

Algorithm 5 Metropolis-Hastings Algorithm

Input: Model P , Proposal Q

- 1: Initialize x_0
- 2: **for** $s = 0, 1, \dots$ **do**
- 3: Sample $x' \sim q(x'|x_s)$
- 4: Compute acceptance probability

$$r = \min \left(1, \frac{p(x')}{p(x_s)} \frac{q(x_s|x')}{q(x'|x_s)} \right)$$

- 5: Sample $u \sim \text{Uniform}(0, 1)$
- 6: **if** $u < r$ **then**
- 7: Accept x'
- 8: Set $x_{s+1} = x'$
- 9: **else**
- 10: Reject x'
- 11: Set $x_{s+1} = x_s$
- 12: **end if**
- 13: **end for**

Output: Accepted Samples

Metropolis-Hastings: Example

Metropolis-Hastings: Conditions

Variants of Metropolis-Hastings

- ▶ Metropolis Algorithm
 - ▶ Uses a symmetric proposal distribution Q , such that:
 $q(x_s|x') = q(x'|x_s)$
- ▶ Metropolis-Adjusted Langevin Algorithm
 - ▶ New states are proposed with Langevin dynamics:
 $x' = x_s + \tau \nabla \log \pi(x_s) + \sqrt{2\tau} \xi_k$
 - ▶ Proposal probabilities are normally distributed:
 $q(x'|x_s) \sim \mathcal{N}(x_s + \tau \nabla \log \pi(x_s), 2\tau I_d)$
- ▶ Hamiltonian Monte-Carlo
 - ▶ Adds Hamiltonian Dynamics
- ▶ No-U-Turn Sampler (NUTS)
 - ▶

Gibbs Sampling

Further Reading

- ▶ Machine Learning: A Probabilistic Perspective by Kevin Murphy
- ▶ Monte Carlo Methods in Financial Engineering by Paul Glasserman
- ▶ Probabilistic Graphical Models: Principles and Techniques by Daphne Koller and Nir Friedman