Introduction to Random Numbers, Sampling, and MCMC Methods

Bill Watson

S&P Global

August 5, 2019

What is Sampling and why is it useful?

- Sampling is the practice of generating observations from a population
- Monte Carlo methods are algorithms that rely on repeated random sampling to obtain approximations where it is difficult or impossible to use deterministic approaches
 - Optimization
 - Numerical Integration
 - Sampling distributions

Application: Approximating π

Algorithm 1 Approximating π

Input: Batch Size N

1: Sample $u_1 \sim \text{Uniform}(0,1) \ N \text{ times}$

2: Sample $u_2 \sim \mathsf{Uniform}(0,1) \ \mathsf{N} \ \mathsf{times}$

3:
$$\tilde{\pi} = \frac{4}{N} \cdot \left| \left\{ (u_1, u_2) \mid \sqrt{u_1^2 + u_2^2} < 1 \right\} \right|$$

Output: $\tilde{\pi}$

Application: Approximating π

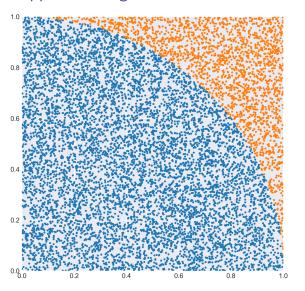
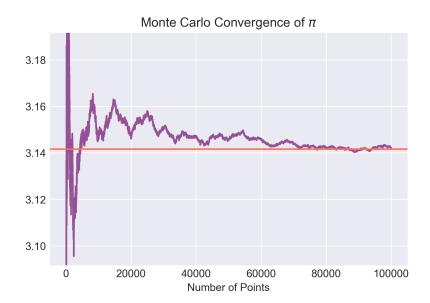


Figure: The ratio of points inside the unit circle approximates $\frac{\pi}{4}$

Application: Approximating π



What can we do with our samples?

- ▶ Integration: $\int p(x) f(x) dx = \frac{1}{n} \sum_{x_i \sim P} f(x_i)$
- Expectation: $\mu = \frac{1}{n} \sum_{x_i} x_i$
- Variance: $\sigma^2 = \frac{1}{n} \sum_{x_i} (x_i \mu)^2$
- ▶ Median: median = median $(x_1, x_2, ... x_n)$
- ► Entropy: $\mathbb{H}(P) = -\frac{1}{n} \sum_{x_i \sim P} \log p(x_i)$
- ► CDF: $p(c) = \frac{1}{n} |\{x_i | x_i \le c\}|$

Pseudo-Random Number Generators

Pseudo-Random Number Generators: LCG

Algorithm 2 Linear Congruential Generator

Input: Modulus m, Multiplier a, Increment c, Seed X_0

1: $X_{i+1} = (a \cdot X_i + c) \mod m$

Output: X_{i+1}

Pseudo-Random Number Generators: Mersenne Twister

Generating a Normal from the CDF

- We can use the cumulative distribution function to sample any distribution
- For instance, a normal's CDF is:

$$F(x) = \frac{1}{2} \left[1 + \operatorname{erf}\left(\frac{x - \mu}{\sigma\sqrt{2}}\right) \right]$$

With an Inverse CDF as:

$$F^{-1}(p) = \mu + \sigma\sqrt{2} \cdot \text{erf}^{-1}(2p-1)$$

erf(x) is the error function, defined as:

$$\operatorname{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$$

Inverse Transform Sampling

It's easy to generalize this method to any distribution with a closed-form inverse CDF

Algorithm 3 Inverse Transform Sampling

Input: Inverse CDF F^{-1}

1: Sample $u \sim \mathsf{Uniform}(0,1)$

2:
$$X = F^{-1}(u)$$

Output: X

Inverse Transform Sampling: Example

Table of Inverse CDFs for Common Distributions

Distribution	$F^{-1}(p)$
$\mathcal{N}(\mu,\sigma)$	$\mu + \sigma\sqrt{2}\cdoterf^{-1}\left(2p-1\right)$
Uniform(a,b)	$a+p\cdot (b-a)$
$Exp(\lambda)$	$rac{-\ln(1- ho)}{\lambda}$

Rejection Sampling

Rejection Sampling: Algorithm

Algorithm 4 Rejection Sampling

Input: Model P, Proposal Q, M > 1

- 1: Sample $x \sim P$
- 2: Sample $u \sim \mathsf{Uniform}(0,1)$
- 3: if $u < \frac{p(x)}{M \cdot q(x)}$ then
- 4: Accept x
- 5: **else**
- 6: Reject x
- 7: end if

Output: Accepted Samples

Rejection Sampling: Example

Rejection Sampling: Pros & Cons

What is MCMC?

Metropolis-Hastings

Metropolis-Hastings: Algorithm

Algorithm 5 Metropolis-Hastings Algorithm

```
Input: Model P, Proposal Q
 1: Initialize xn
 2: for s = 0, 1, \dots do
     Sample x' \sim q(x'|x_s)
        Compute acceptance probability
 4:
        r = \min\left(1, \frac{p(x')}{p(x_s)} \frac{q(x_s|x')}{q(x'|x_s)}\right)
        Sample u \sim \text{Uniform}(0,1)
 5:
        if u < r then
 6.
           Accept x'
 7.
          Set x_{s+1} = x'
 8:
       else
 g.
           Reject x'
10:
           Set x_{s+1} = x_s
11:
        end if
12.
13: end for
Output: Accepted Samples
```

Metropolis-Hastings: Example

Metropolis-Hastings: Conditions

Variants of Metropolis-Hastings

- Metropolis Algorithm
 - ▶ Uses a symmetric proposal distribution Q, such that: $q(x_s|x') = q(x'|x_s)$
- Metropolis-Adjusted Langevin Algorithm
 - ▶ New states are proposed with Langevin dynamics: $x' = x_s + \tau \nabla \log \pi(x_s) + \sqrt{2\tau} \xi_k$
 - Proposal probabilties are normally distributed: $q(x'|x_s) \sim \mathcal{N}(x_s + \tau \nabla \log \pi(x_s), 2\tau I_d)$
- Hamiltonian Monte-Carlo
- No-U-Turn Sampler (NUTS)

Gibbs Sampling

Further Reading

- ► Machine Learning: A Probabilistic Perspective by Kevin Murphy
- Monte Carlo Methods in Financial Engineering by Paul Glasserman
- ► Probabilistic Graphical Models: Principles and Techniques by Daphne Koller and Nir Friedman